

Stat 102C HW5: Answer Key

Jonathan Arfa, Muzhou Liang

June 6, 2014

Problem 1.1

$$P(X_{t+1} = x, Y_t = y) = P(Y_t = y)P(X_{t+1} = x|Y_t = y) \quad (1)$$

$$= p_Y(y)p(x|y) \quad (2)$$

$$= p(x, y) \quad (3)$$

$$P(X_{t+1} = x, Y_{t+1} = y) = P(X_{t+1} = x)P(Y_{t+1} = y|X_{t+1} = x) \quad (4)$$

$$= p_X(x)p(y|x) \quad (5)$$

$$= p(x, y) \quad (6)$$

Thus, $(X_{t+1}, Y_t) \sim p(x, y)$ and $(X_{t+1}, Y_{t+1}) \sim p(x, y)$.

Problem 1.2

```
(probs = matrix(c(0.1, 0.4, 0.3, 0.2), nrow = 2, byrow = TRUE))

##      [,1] [,2]
## [1,]  0.1  0.4
## [2,]  0.3  0.2

# marginal probabilities of x: p(x==1), p(x==2)
(px = rowSums(probs))

## [1] 0.5 0.5

# marginal probabilities of y
(py = colSums(probs))

## [1] 0.4 0.6

# conditional probabilities of x/y==1: p(x==1/y==1), p(x==2/y==1)
(px_y1 = c(probs[1, 1], probs[2, 1])/sum(probs[, 1]))
```

```
## [1] 0.25 0.75

# conditional probabilities of x/y==2
(px_y2 = c(probs[1, 2], probs[2, 2])/sum(probs[, 2]))

## [1] 0.6667 0.3333

# conditional probabilities of y/x==1
(py_x1 = c(probs[1, 1], probs[1, 2])/sum(probs[1, ]))

## [1] 0.2 0.8

# conditional probabilities of y/x==2
(py_x2 = c(probs[2, 1], probs[2, 2])/sum(probs[2, ]))

## [1] 0.6 0.4
```

For the last part of this problem, we need to show that
 $p(x, y) = P(X_{t+1} = x | Y_t = y) p_Y(y)$.
 So, $p(x = 1, y = 1) = P(X_{t+1} = 1 | Y_t = 1) p_Y(1)$.

```
probs[1, 1]

## [1] 0.1

px_y1[1] * py[1]

## [1] 0.1
```

$$p(x = 1, y = 2) = P(X_{t+1} = 1 | Y_t = 2) p_Y(2).$$

```
probs[1, 2]

## [1] 0.4

px_y2[1] * py[2]

## [1] 0.4
```

$$p(x = 2, y = 1) = P(X_{t+1} = 2 | Y_t = 1) p_Y(1).$$

```
probs[2, 1]

## [1] 0.3

px_y1[2] * py[1]

## [1] 0.3
```

$$p(x = 2, y = 2) = P(X_{t+1} = 2 | Y_t = 2) p_Y(2).$$

```
probs[2, 2]
## [1] 0.2
px_y2[2] * py[2]
## [1] 0.2
```

Problem 1.3

```
n = 5
# initialize a 3-dimensional array
jointdist = array(dim = c(2, 2, n + 1))
# set the first matrix of that array to p0
jointdist[, , 1] <- matrix(c(1, 0, 0, 0), nrow = 2)
for (t in 2:(n + 1)) {
  # get x_t/y_(t-1)
  x1y1 = px_y1[1] * sum(jointdist[, 1, t - 1])
  x2y1 = px_y1[2] * sum(jointdist[, 1, t - 1])
  x1y2 = px_y2[1] * sum(jointdist[, 2, t - 1])
  x2y2 = px_y2[2] * sum(jointdist[, 2, t - 1])
  jointdist[, , t] = matrix(c(x1y1, x1y2, x2y1, x2y2), 2, byrow = TRUE)

  # get y_t/x_t
  x1y1 = py_x1[1] * sum(jointdist[1, , t])
  x1y2 = py_x1[2] * sum(jointdist[1, , t])
  x2y1 = py_x2[1] * sum(jointdist[2, , t])
  x2y2 = py_x2[2] * sum(jointdist[2, , t])
  jointdist[, , t] = matrix(c(x1y1, x1y2, x2y1, x2y2), 2, byrow = TRUE)
}
for (t in 0:n) {
  print(paste("Joint distribution at t =", t))
  print(jointdist[, , t + 1])
}


## [1] "Joint distribution at t = 0"
##      [,1] [,2]
## [1,]    1    0
## [2,]    0    0
## [1] "Joint distribution at t = 1"
##      [,1] [,2]
## [1,] 0.05 0.2
```

```
## [2,] 0.45 0.3
## [1] "Joint distribution at t = 2"
##      [,1] [,2]
## [1,] 0.09167 0.3667
## [2,] 0.32500 0.2167
## [1] "Joint distribution at t = 3"
##      [,1] [,2]
## [1,] 0.09861 0.3944
## [2,] 0.30417 0.2028
## [1] "Joint distribution at t = 4"
##      [,1] [,2]
## [1,] 0.09977 0.3991
## [2,] 0.30069 0.2005
## [1] "Joint distribution at t = 5"
##      [,1] [,2]
## [1,] 0.09996 0.3998
## [2,] 0.30012 0.2001
```

This isn't required by the homework, but we can create a metric of how close our approximate joint distribution is to p at each value of t .

```
distance = apply(jointdist, 3, function(mat) {
  # the distance between two matrices is the sqrt of the sum of the squared
  # differences
  sqrt(sum((mat - probs)^2))
})
plot(0:n, distance, xlab = "t", ylab = "", type = "o", pch = 19, main = "Distance from p")
```

Distance from p



Problem 1.4

```
Nsim = 20
M = 1e+05
# create a data frame of each person's x_t y value
distribution = data.frame(x = rep(1, M), y = rep(1, M))
jointdist = array(0, dim = c(2, 2, Nsim))
for (t in 1:Nsim) {
  # get (x_t/y_{t-1})
  r = runif(M) #generate a random number for each person
  indy1 = which(distribution$y == 1) #indices of which people have y_{t-1}=1
  indy2 = which(distribution$y == 2)
  distribution$x[indy1] = ifelse(r[indy1] < px_y1[1], 1, 2)
  distribution$x[indy2] = ifelse(r[indy2] < px_y2[1], 1, 2)

  # get (y_t/x_t)
  r = runif(M)
  indx1 = which(distribution$x == 1)
  indx2 = which(distribution$x == 2)
```

```

distribution$y[indx1] = ifelse(r[indx1] < py_x1[1], 1, 2)
distribution$y[indx2] = ifelse(r[indx2] < py_x2[1], 1, 2)
for (x in 1:2) for (y in 1:2) {
  # fill in each cell of the joint dist.
  jointdist[x, y, t] = sum(distribution$x == x & distribution$y == y)
}
# convert from number of people to proportion
jointdist[, , t] = jointdist[, , t]/M
}
jointdist[, , Nsim]

##          [,1]    [,2]
## [1,] 0.09889 0.3968
## [2,] 0.30230 0.2021


probs

##          [,1] [,2]
## [1,]    0.1  0.4
## [2,]    0.3  0.2

distance = apply(jointdist, 3, function(mat) {
  sqrt(sum((mat - probs)^2))
})
plot(1:Nsim, distance, xlab = "t", ylab = "", type = "o", pch = 19, main = paste0("Distance
M))
axis(1, at = c(1, seq(5, Nsim, by = 5)))

```

Distance from p with M = 1e+05



Problem 2

Since we want to test out various values of ρ , let's create a function that we can run with different parameters.

```
gibbsBivarNorm = function(rho, Nsim = 2000, burnin = 1000, x0 = 0, y0 = 0, returnvals = TRUE,
  plotit = TRUE) {
  x = y = numeric(Nsim)
  # first iteration happens outside the for loop
  x[1] = rnorm(1, rho * y0, sqrt(1 - rho^2))
  y[1] = rnorm(1, rho * x[1], sqrt(1 - rho^2))
  for (t in 2:Nsim) {
    x[t] = rnorm(1, rho * y[t - 1], sqrt(1 - rho^2)) #conditional dist. of x/y
    y[t] = rnorm(1, rho * x[t], sqrt(1 - rho^2)) #conditional dist. of y/x
  }
  # remove the samples from the burn-in period.
  x = x[(burnin + 1):Nsim]
  y = y[(burnin + 1):Nsim]
  if (plotit)
    plot(x, y, main = paste("rho =", rho), xlim = c(-4, 4), ylim = c(-4,
```

```

4))
# note that I'm keeping the limits of the graph constant regardless of rho
if (returnvals)
  return(list(x = x, y = y, expec_val = mean(sqrt(x^2 + y^2))))
}
par(mfrow = c(2, 2), mar = c(2.5, 2, 2, 1) + 0.1)
# FYI, default margins are mar=c(5, 4, 4, 2) + 0.1
gibbsBivarNorm(-0.9)$expec_val

## [1] 1.177

gibbsBivarNorm(-0.3)$expec_val

## [1] 1.254


gibbsBivarNorm(0.1)$expec_val

## [1] 1.276

gibbsBivarNorm(0.9)$expec_val

## [1] 1.16

```



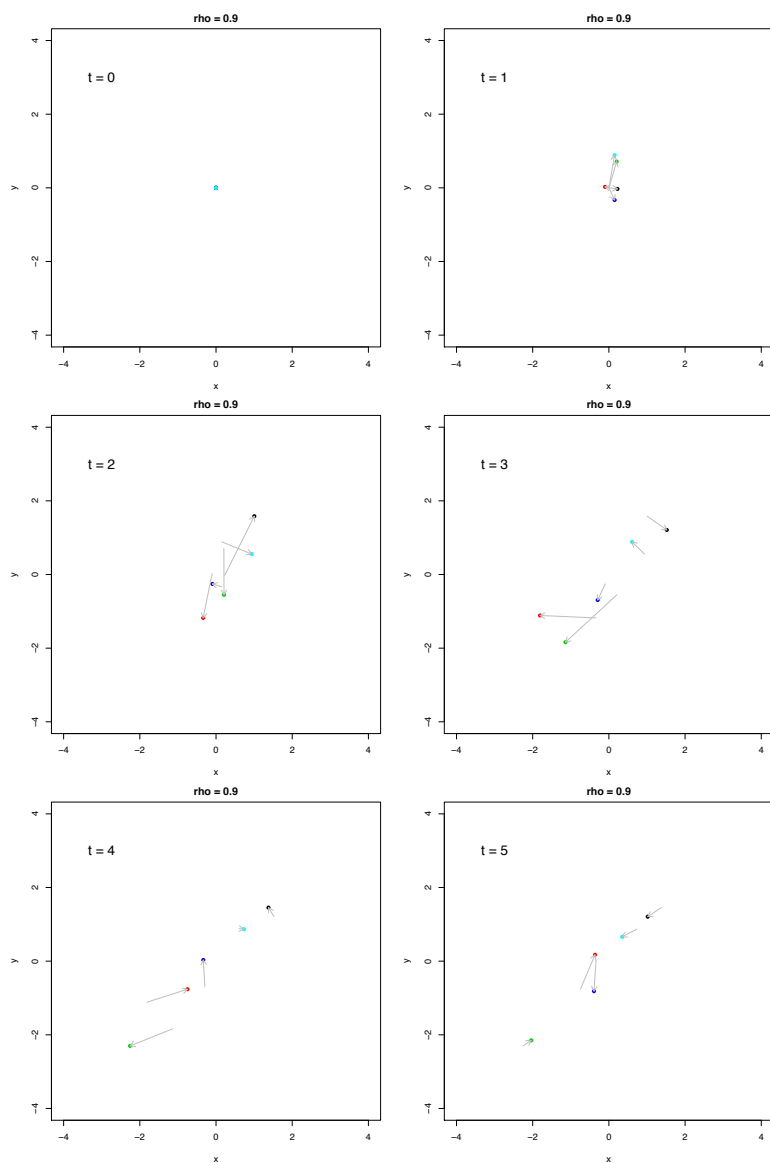
Problem 2-Bonus

We'll create a function that outputs a series of plots, interrupted by pauses. If we wrap this function in the `saveGIF()` function it should create an animated gif.

```
library(animation)
gibbsBivarNorm_walk = function(rho, M, x0 = 0, y0 = 0, Nsim = 50, burnin = 0,
  arrow = TRUE) {
  xvals = yvals = matrix(nrow = Nsim + 1, ncol = M)
  for (i in 1:M) {
    # get the samples from each walk separately
    samples = gibbsBivarNorm(rho, Nsim, burnin, x0, y0, plotit = FALSE)
    xvals[, i] = c(x0, samples$x)
    yvals[, i] = c(y0, samples$y)
  }
  for (t in (burnin + 1):(Nsim + 1)) {
    # plot all samples from a given t together
    plot(xvals[t, ], yvals[t, ], type = "p", col = 1:M, xlim = c(-4, 4),
      ylim = c(-4, 4), xlab = "x", ylab = "y", pch = 16, main = paste("rho =",
        rho))
    # draw arrows from previous points
    if (arrow & t > 1)
      arrows(xvals[t - 1, ], yvals[t - 1, ], xvals[t, ], yvals[t, ], length = 0.1,
        col = "grey")
    # write t (notice that I'm placing the text where it's less likely to
    # overlap text)
    if (rho >= 0) {
      text(-3, 3, paste("t =", t - 1), cex = 1.5)
    } else text(3, 3, paste("t =", t - 1), cex = 1.5)
    ani.pause()
  }
}
```

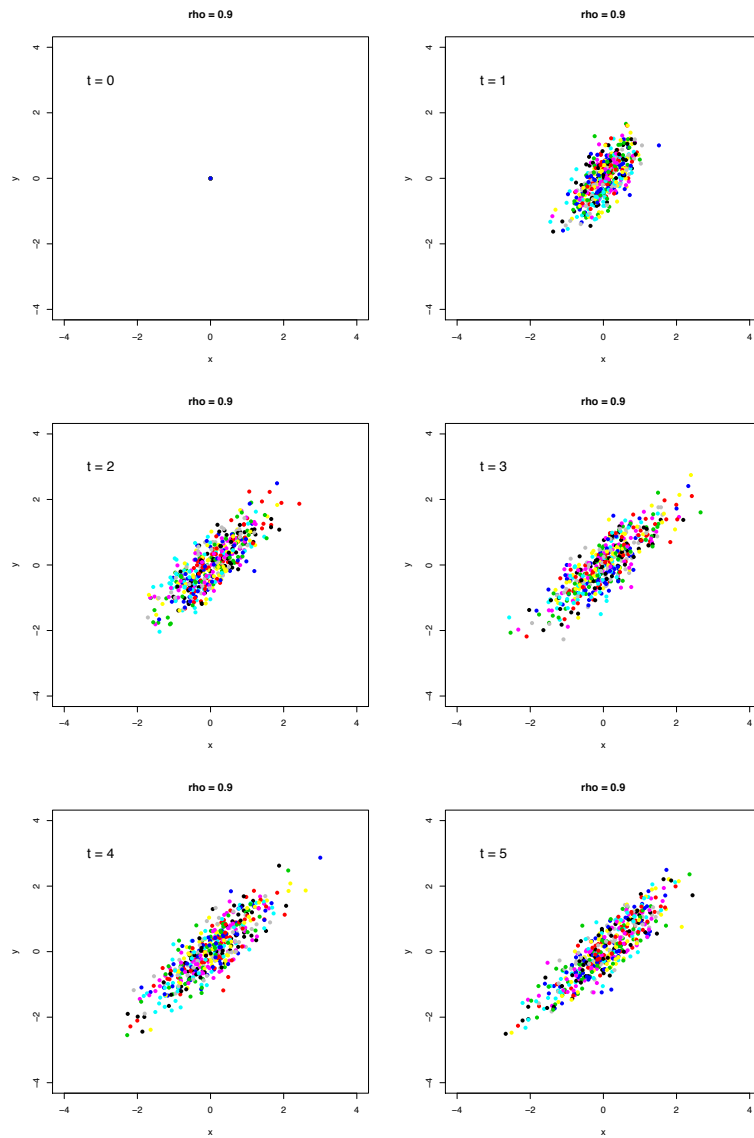
Let's check out a few of these visualizations. Here we're using a low M, starting all of our chains at (0,0), and using arrows to see how each point is moving.

```
par(mfrow = c(1, 1), mar = c(4, 4, 2, 1))
ani.options(interval = 0.1)
# the next 2 lines will generate the 'video' in your plotting window
gibbsBivarNorm_walk(rho = 0.9, M = 5, Nsim = 5)
```



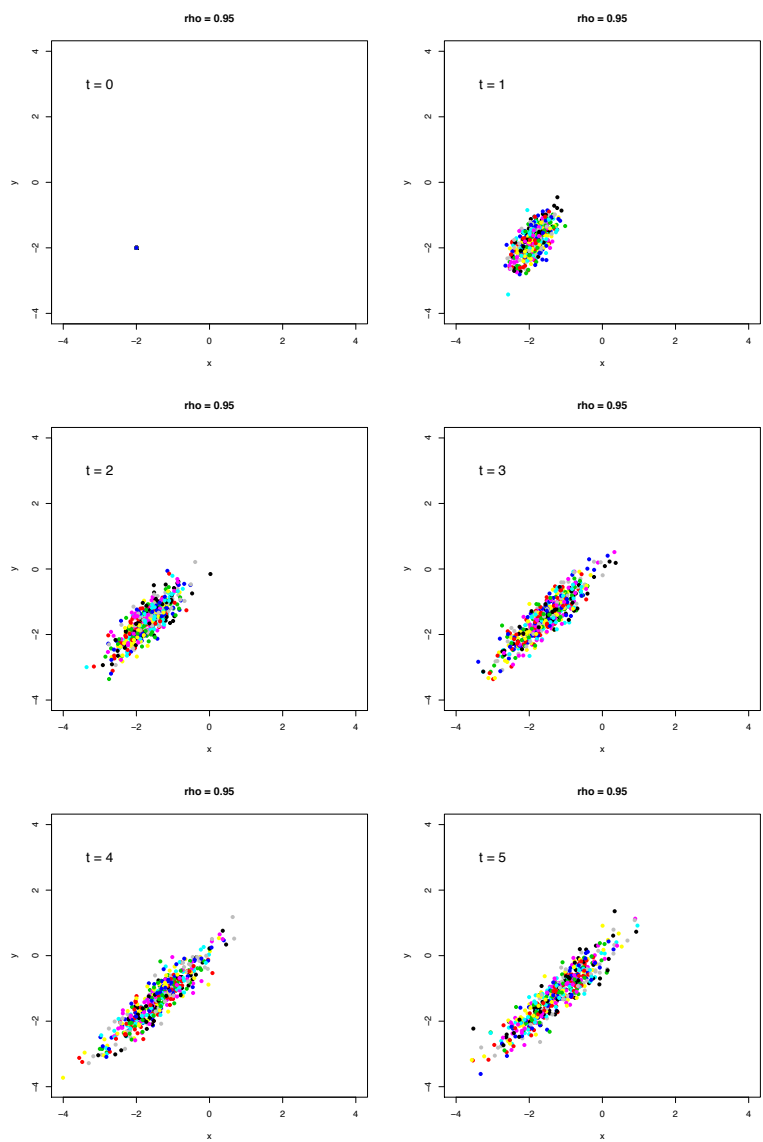
Here's a much larger M. See how the distribution is covering to the joint distribution $p(x, y)$?


```
gibbsBivarNorm_walk(rho = 0.9, M = 500, arrow = FALSE, Nsim = 5)
```



And finally, here are 500 Markov chains that are starting far from the center and with $\rho = 0.95$.

```
gibbsBivarNorm_walk(rho = 0.95, M = 500, arrow = FALSE, x0 = -2, y0 = -2, Nsim = 9)
```





The following code chunk isn't being run, but will save GIFs of the sequences.

```
# The file should pop up separately, make sure you save it somewhere useful.
saveGIF(gibbsBivarNorm_walk(rho = 0.9, M = 5, Nsim = 50), movie.name = "smallM.gif")
saveGIF(gibbsBivarNorm_walk(rho = 0.9, M = 500, arrow = FALSE, Nsim = 50), movie.name = "bigM1.gif")
saveGIF(gibbsBivarNorm_walk(rho = 0.95, M = 500, arrow = FALSE, x0 = -3, y0 = -3,
                             Nsim = 50), movie.name = "bigM2.gif")
```